

Lecture 18 R scripts

Example: Variational Bayes applied to farm data.

To illustrate Variational Bayes, we look at an agricultural dataset, which can be downloaded as `farmdata.csv` (for weaning weights) and `Kmat.csv` (for the kinship matrix) from Canvas. To benchmark Variational Bayes, we will compare results to those obtained using Gibbs sampling.

Writing a Gibbs sampler for this problem

The Gibbs sampler for this problem is just a modification of the Gibbs sampler written in lab 7. Now however we can work with arbitrary variance-covariance matrices in the normal prior by including the argument `Kinv`. $\text{Kinv} = \mathbf{K}^{-1}$ is the inverse of the prior variance-covariance matrix up to proportionality.

```
#Arguments are
#iter: no of iterations
#Z: Predictor matrix for random effects
#X: Predictor matrix for fixed effects
#y: response vector
#burnin: number of initial iterations to discard.
#taue_0: initial guess for residual precision.
#tauu_0: initial guess for random effect precision
#Kinv: inverse of K, where p(u) = N(0, \sigma^2_u K)
#a.u, b.u: hyper-parameters of gamma prior for tauu
#a.e, b.e: hyper-parameters of gamma prior for taue

normalmm.Gibbs<-function(iter,Z,X,y,burnin,taue_0,tauu_0,Kinv,a.u,b.u,a.e,b.e){

n  <-length(y) #no. observations
p  <-dim(X)[2] #no of fixed effect predictors.
q  <-dim(Z)[2] #no of random effect levels
tauu<-tauu_0
taue<-taue_0
beta0<-rnorm(p)
u0   <-rnorm(q,0,sd=1/sqrt(tauu))

#Building combined predictor matrix.
W<-cbind(X,Z)
WTW <-crossprod(W)
WTy <-crossprod(W,y)
library(mvtnorm)

#storing results.
par <-matrix(0,iter,p+q+2)
#Calculating log predictive densities
lppd<-matrix(0,iter,n)

#Create modified identity matrix for joint posterior.
I0  <-diag(p+q)
```

```

diag(I0)[1:p]<-0
I0[-c(1:p),-c(1:p)] <-Kinv

for(i in 1:iter){
#Conditional posteriors.
  uKinvu <- t(u0) %*% Kinv %*% u0
  uKinvu <-as.numeric(uKinvu)
  tauu <-rgamma(1,a.u+0.5*q,b.u+0.5*uKinvu)
#Updating component of normal posterior for beta,u
  Prec <-WTW + tauu*I0/taue
  P.mean<- solve(Prec)%*%WTy
  P.var <-solve(Prec)/taue
  betau <-rmvnorm(1,mean=P.mean,sigma=P.var)
  betau <-as.numeric(betau)
  err <- y-W%*%betau
  taue <-rgamma(1,a.e+0.5*n,b.e+0.5*sum(err^2))
#storing iterations.
  par[i,]<-c(betau,1/sqrt(tauu),1/sqrt(taue))
  beta0 <-betau[1:p]
  u0 <-betau[p+1:q]
  lppd[i,]= dnorm(y,mean=as.numeric(W%*%betau),sd=1/sqrt(taue))
}

lppd      = lppd[-c(1:burnin),]
lppdest   = sum(log(colMeans(lppd)))           #Estimating lppd for whole dataset.
pwaic2    = sum(apply(log(lppd),2,FUN=var)) #Estimating effective number of parameters.
par <-par[-c(1:burnin),]
colnames(par)<-c(paste('beta',1:p,sep=''),paste('u',1:q,sep=''),'sigma_b','sigma_e')
mresult<-list(par,lppdest,pwaic2)
names(mresult)<-c('par','lppd','pwaic')
return(mresult)
}

```

Importing and formatting the data

```

#data<-read.csv('farmdata.csv')
data<-read.csv(file.choose()) #Choose farmdata.csv as downloaded from Canvas
#K <-read.csv('Kmat.csv')
K<-read.csv(file.choose()) #Choose farmdata.csv as downloaded from Canvas
K <-as.matrix(K)
Kinv<-solve(K)

n<-dim(data)[1]
q<-dim(Kinv)[1]
X<-table(1:n,data$flock) #flock is fixed effect
#Indicator matrix for parents.
Z2<-table(1:n,data$sire)
Z3<-cbind(Z2,table(1:n,data$dam))

```

running the Gibbs sampler, checking convergence, and calculating effective sample size

```
system.time(chain1<-normalmm.Gibbs(iter=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=5,tauu_0=0.2,Kinv=Kinv)
```

```

## Warning: package 'mvtnorm' was built under R version 4.0.5
##       user  system elapsed
##      5.08    0.12   5.33

system.time(chain2<-normalmm.Gibbs(iter=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=1,tauu_0=1,Kinv=Kinv)

##       user  system elapsed
##      5.05    0.09   5.22

system.time(chain3<-normalmm.Gibbs(iter=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=0.2,tauu_0=3,Kinv=Kinv)

##       user  system elapsed
##      5.08    0.14   5.27

library(coda)

## Warning: package 'coda' was built under R version 4.0.5

rml1<-as.mcmc.list(as.mcmc((chain1$par[1:4000])))
rml2<-as.mcmc.list(as.mcmc((chain2$par[1:4000])))
rml3<-as.mcmc.list(as.mcmc((chain3$par[1:4000])))
rml4<-as.mcmc.list(as.mcmc((chain1$par[4000+1:4000])))
rml5<-as.mcmc.list(as.mcmc((chain2$par[4000+1:4000])))
rml6<-as.mcmc.list(as.mcmc((chain3$par[4000+1:4000])))
rml<-c(rml1,rml2,rml3,rml4,rml5,rml6)

#Gelman-Rubin diagnostic.
gelman.diag(rml)[[1]]
```

	Point est.	Upper C.I.
## beta1	1.0001886	1.000618
## beta2	1.0003186	1.000664
## u1	1.0002849	1.000427
## u2	1.0004086	1.000672
## u3	1.0002609	1.000783
## u4	1.0002456	1.000754
## u5	1.0001370	1.000347
## u6	1.0001749	1.000684
## u7	1.0001805	1.000262
## u8	1.0000731	1.000364
## u9	1.0001241	1.000346
## u10	1.0000389	1.000371
## u11	0.9999831	1.000121
## u12	1.0000541	1.000331
## u13	1.0000604	1.000143
## sigma_b	1.0009851	1.002606
## sigma_e	1.0012429	1.002924

```
#effective sample size.
effectiveSize(rml)
```

	beta1	beta2	u1	u2	u3	u4	u5	u6
##	24000.000	24000.000	24000.000	21692.020	23909.883	24000.000	23848.486	24000.000
##	u7	u8	u9	u10	u11	u12	u13	sigma_b
##	24000.000	25176.790	23691.012	23376.080	24000.000	24000.000	24000.000	10917.902
##	sigma_e							
##	8676.271							

Writing E-M algorithm to find V-B estimates

To determine the parameters of the Variational Bayes approximation, we will use an Expectation-Maximisation algorithm. This means we will successively update $E(\beta)$, $\text{Var}(\beta)$, $E(\mathbf{u})$, $\text{Var}(\mathbf{u})$, τ_u , τ_e until convergence has been reached. In this code, convergence is defined as when $\sqrt{(\theta^{(t)} - \theta^{(t-1)})^2 / (|\theta^{(t)}| + 0.01)} \leq \epsilon$ for all elements of θ . The addition of 0.01 is to avoid problems when $\theta^{(t)} = 0$

```
#Arguments are
#epsilon: accuracy cut-off.
#iter: no of iterations
#Kinv: inverse of K, where p(u) = N(0, \sigma^2_u K)
#Z: Predictor matrix for random effects
#X: Predictor matrix for fixed effects
#y: response vector
#taue_0: initial guess for residual precision.
#tauu_0: initial guess for random effect precision
#a.u, g.u: hyper-parameters of gamma prior for tauu
#a.e, g.e: hyper-parameters of gamma prior for taue

#Output are final estimates, plus iteration number when convergence was reached.
VB.mm<-function(epsilon,iter,Kinv,Z,X,y,taue_0,tauu_0,u0,beta0,a.e,g.e,a.u,g.u){
  n<-dim(X)[1]
  p<-dim(X)[2]
  q<-dim(Z)[2]
  ZTZ<-crossprod(Z)
  ZTY<-crossprod(Z,y)
  XTY<-crossprod(X,y)
  ZTX<-crossprod(Z,X)
  XTX<-crossprod(X)
  XTXinv<-solve(XTX)

  for(i in 1:iter){
    Vu <-solve(taue_0*ZTZ+tauu_0*Kinv) #update Var(u)
    u <-taue_0*Vu%*%(ZTY-ZTX%*%beta0) #update E(u)
    Vb <-XTXinv/taue_0 #update Var(beta)
    b <-XTXinv%*%(XTY-t(ZTX)%*%u) #update E(beta)
    TrKinu <- sum(diag(Kinv%*%Vu))
    uKinu <- t(u)%*%Kinu%*%u
    tauu <- (a.u+0.5*q)/(g.u+0.5*as.numeric(uKinu)+0.5*TrKinu)
    tauu <- as.numeric(tauu)
    err <- y - X%*%b - Z%*%u
    TrXTXb <- sum(diag(XTX%*%Vb))
    TrZTZu <- sum(diag(ZTZ%*%Vu))
    taue <- (a.e+0.5*n)/(g.e+0.5*sum(err^2)+0.5*TrXTXb+0.5*TrZTZu)
    taue <- as.numeric(taue)

    if(i > 1){
      diffb <- sqrt((b-beta0)^2)/(abs(b)+0.01)
      diffu <- sqrt((u-u0)^2)/(abs(u)+0.01)
      diffte <- abs(taue_0-taue)/(taue+0.01)
      difftu <- abs(tauu_0-tauu)/(tauu+0.01)
      diffvb <- sqrt((diag(Vb0) - diag(Vb))^2)/(diag(Vb))
      diffvu <- sqrt((diag(Vu0) - diag(Vu))^2)/(diag(Vu))
      diff.all<-c(diffb,diffu,diffte,difftu,diffvb,diffvu)
      if(max(diff.all) < epsilon) break
    }
  }
}
```

```

}
Vu0 <- Vu;u0<-u;Vb0<-Vb;beta0<-b;taue_0<-taue;tauu_0<-tauu
#Calculate relative change.
}

taue.param<-c((a.e+0.5*n),(g.e+0.5*sum(err^2)+0.5*TrXTXb+0.5*TrZTZu))
tauu.param<-c((a.u+0.5*q),(g.u+0.5*uKinvu+0.5*TrKinvu))
param<-list(b,Vb,u,Vu,taue.param,tauu.param,i)
names(param)<-c('beta_mean','beta_var','u_mean','u_var','tau_e','tau_u','iter')
return(param)
}

```

Estimating parameter of variational Bayes approximations

```

system.time(test1<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=0.2,tauu_0=0.2,u0=rn
##      user    system elapsed
##      0.29      0.00     0.29

system.time(test2<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=5,tauu_0=1,u0=rnorm(
##      user    system elapsed
##      0.19      0.00     0.19

system.time(test3<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=1,tauu_0=5,u0=rnorm(
##      user    system elapsed
##      0.14      0.00     0.14

test1$iter
## [1] 1070
test2$iter
## [1] 1152
test3$iter
## [1] 870

```

Comparing estimates obtained from Gibbs sampling and Variational Bayes

```

#Comparing point estimates/posterior means

chain.all<-rbind(chain1$par,chain2$par,chain3$par)
#beta
cbind(test1$beta_mean,test2$beta_mean,test3$beta_mean,colMeans(chain.all[,1:2]))

##      [,1]      [,2]      [,3]      [,4]
## 1 4.149860 4.149861 4.149861 4.153120
## 2 4.708765 4.708766 4.708767 4.709322

#u
cbind(test1$u_mean,test2$u_mean,test3$u_mean,colMeans(chain.all[,3:15]))

##      [,1]      [,2]      [,3]      [,4]
## 1 0.80578542 0.80578469 0.80578421 0.80841251

```

```

## 2 -1.08621855 -1.08621928 -1.08621975 -1.09169227
## 3 0.21245313 0.21245245 0.21245200 0.21098242
## 4 0.47527965 0.47527943 0.47527928 0.47601605
## 5 -0.52479170 -0.52479204 -0.52479226 -0.53454094
## 6 -0.69870349 -0.69870374 -0.69870390 -0.70499465
## 7 0.50963975 0.50963940 0.50963918 0.51611965
## 8 -0.16302215 -0.16302238 -0.16302253 -0.16250230
## 9 -0.15022031 -0.15022065 -0.15022087 -0.14939405
## 10 0.12955976 0.12955948 0.12955930 0.13095244
## 11 0.14797212 0.14797175 0.14797150 0.15130713
## 12 -0.07160937 -0.07160964 -0.07160982 -0.07495579
## 13 0.57213111 0.57213069 0.57213042 0.57072883

#Variances.
sigmae2<-sigmau2<-rep(0,3)
sigmae2[1]<-test1$tau_e[2]/(test1$tau_e[1]-1)
sigmau2[1]<-test1$tau_u[2]/(test1$tau_u[1]-1)
sigmae2[2]<-test2$tau_e[2]/(test2$tau_e[1]-1)
sigmau2[2]<-test2$tau_u[2]/(test2$tau_u[1]-1)
sigmae2[3]<-test3$tau_e[2]/(test3$tau_e[1]-1)
sigmau2[3]<-test3$tau_u[2]/(test3$tau_u[1]-1)
cbind(sigmau2,mean(chain.all[,16]^2))

## sigmau2
## [1,] 0.352975 0.4649196
## [2,] 0.352975 0.4649196
## [3,] 0.352975 0.4649196
cbind(sigmae2,mean(chain.all[,17]^2))

## sigmae2
## [1,] 0.05642792 0.05242605
## [2,] 0.05642792 0.05242605
## [3,] 0.05642792 0.05242605

#Comparing posterior distributions.
par(mfrow=c(5,4))
mlim<-quantile(chain.all[,1],c(0.005,0.995))
curve(dnorm(x,mean=test1$beta_mean[1],sd=sqrt(test1$beta_var[1,1])),ylab='Density',main='',xlim=mlim,co
lines(density(chain.all[,1]))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)

mlim<-quantile(chain.all[,2],c(0.005,0.995))
curve(dnorm(x,mean=test1$beta_mean[2],sd=sqrt(test1$beta_var[2,2])),ylab='Density',main='',xlim=mlim,co
lines(density(chain.all[,2]))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)

#Repeat for random effects.
for(i in 1:13){
  mlim<-quantile(chain.all[,i+2],c(0.005,0.995))
  curve(dnorm(x,mean=test1$u_mean[i],sd=sqrt(test1$u_var[i,i])),ylab='Density',main='',xlim=mlim,col=2,lty=1,bty='n',cex=2.5)
  lines(density(chain.all[,i+2]))
  legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)
}

```

```

mlim<-quantile(chain.all[,16]^2,c(0.005,0.995))
curve(dgamma(1/x,shape=test1$tau_u[1],rate=test1$tau_u[2])*x^(-2),ylab='Density',main='',xlim=c(0,mlim[2])
lines(density(chain.all[,16]^2))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)

mlim<-quantile(chain.all[,17]^2,c(0.005,0.995))
curve(dgamma(1/x,shape=test1$tau_e[1],rate=test1$tau_e[2])*x^(-2),ylab='Density',main='',xlim=c(0,mlim[2])
lines(density(chain.all[,17]^2))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)

```

