

# Lab 9 solutions MAST90125: Bayesian Statistical learning

## Variational Bayes.

In lecture 18, we looked at (mean-field) Variational Bayes as a method to find approximate posterior distributions for sub-blocks of the parameter vector  $\boldsymbol{\theta}$ , for the model

$$\begin{aligned} p(\mathbf{y}|\boldsymbol{\beta}, \mathbf{u}, \tau_e) &= \mathcal{N}(\mathbf{X}\boldsymbol{\beta} + \mathbf{Z}\mathbf{u}, \frac{1}{\tau_e}\mathbf{I}_n) \\ p(\boldsymbol{\beta}) &\propto 1 \\ p(\mathbf{u}) &= \mathcal{N}(\mathbf{0}_q, \frac{1}{\tau_u}\mathbf{K}) \\ p(\tau_e) &= \text{Ga}(\alpha_e, \gamma_e) \\ p(\tau_u) &= \text{Ga}(\alpha_u, \gamma_u) \end{aligned}$$

## Instructions for lab

Download `farmpdata.csv` and `Kmat.csv` from LMS.

Re-format the code given in lecture 18 for performing Variational Bayes so that rather than constructing approximate posteriors,

- $Q(\boldsymbol{\beta})$
- $Q(\mathbf{u})$
- $Q(\tau_e)$
- $Q(\tau_u)$

you determine the approximate posteriors

- $Q(\boldsymbol{\beta}, \mathbf{u})$
- $Q(\tau_e)$
- $Q(\tau_u)$ .

Compare the performance of the new blocking structure to that used in class.

## Solution

The difference between the code given in class and this lab is that we approximate the posterior of  $\boldsymbol{\beta}, \mathbf{u}$  jointly. To do this, note the trick that the joint kernel of  $\boldsymbol{\beta}, \mathbf{u}$  can be written as

$$p(\boldsymbol{\beta}, \mathbf{u}) \propto e^{-\frac{\tau_u(\boldsymbol{\beta}' - \mathbf{u}') \begin{pmatrix} \mathbf{0}_{p \times p} & \mathbf{0}_{p \times q} \\ \mathbf{0}_{q \times p} & \mathbf{K}^{-1} \end{pmatrix} (\boldsymbol{\beta} - \mathbf{u})}{2}},$$

Hence in terms of coding, updating  $\beta, \mathbf{u}$  jointly looks just like updating  $\mathbf{u}$  in the code given in lecture 18, except that  $\beta = \mathbf{0}$ .

As a reminder, the joint distribution,  $p(y, \beta, \mathbf{u}, \tau_e, \tau_u | \mathbf{X}, \mathbf{Z})$  is

$$\left(\frac{\tau_e}{2\pi}\right)^{\frac{n}{2}} e^{-\frac{\tau_e(\mathbf{y}-\mathbf{X}\beta-\mathbf{Z}\mathbf{u})'(\mathbf{y}-\mathbf{X}\beta-\mathbf{Z}\mathbf{u})}{2}} \times 1 \times \left(\frac{\tau_u}{2\pi}\right)^{\frac{q}{2}} \det(\mathbf{K})^{-1/2} e^{-\frac{\tau_u \mathbf{u}' \mathbf{K}^{-1} \mathbf{u}}{2}} \times \frac{\gamma_u^{\alpha_u} \tau_u^{\alpha_u-1} e^{-\gamma_u \tau_u}}{\Gamma(\alpha_u)} \times \frac{\gamma_e^{\alpha_e} \tau_e^{\alpha_e-1} e^{-\gamma_e \tau_e}}{\Gamma(\alpha_e)},$$

and that we need to determine the expectation of log-kernels.

- For  $\tau_e$ , the kernel and log-kernel are respectively

$$\begin{aligned} \text{Kernel: } & \tau_e^{\frac{n}{2}} e^{-\frac{\tau_e(\mathbf{y}-\mathbf{X}\beta-\mathbf{Z}\mathbf{u})'(\mathbf{y}-\mathbf{X}\beta-\mathbf{Z}\mathbf{u})}{2}} \tau_e^{\alpha_e-1} e^{-\gamma_e \tau_e} \\ \text{Log-kernel: } & (n/2 + \alpha_e - 1) \log(\tau_e) - \tau_e(\gamma_e + (\mathbf{y} - \mathbf{X}\beta - \mathbf{Z}\mathbf{u})'(\mathbf{y} - \mathbf{X}\beta - \mathbf{Z}\mathbf{u})/2) \end{aligned}$$

- The expected log-kernel  $E_{-\tau_e}$  (Log-kernel) is

$$(\frac{n}{2} + \alpha_e - 1) \log(\tau_e) - \tau_e(\gamma_e + \frac{\mathbf{y}'\mathbf{y} + E_{\beta,\mathbf{u}} \left( (\beta' \quad \mathbf{u}') \begin{pmatrix} \mathbf{X}'\mathbf{X} & \mathbf{X}'\mathbf{Z} \\ \mathbf{Z}'\mathbf{X} & \mathbf{Z}'\mathbf{Z} \end{pmatrix} \begin{pmatrix} \beta \\ \mathbf{u} \end{pmatrix} \right)}{2} - \mathbf{y}'\mathbf{X}E_{\beta,\mathbf{u}}(\beta) - \mathbf{y}'\mathbf{Z}E_{\beta,\mathbf{u}}(\mathbf{u}))$$

- Using the tricks that  $E(\mathbf{x}\mathbf{x}') = \text{Var}(\mathbf{x}) + E(\mathbf{x})E(\mathbf{x})'$  and that  $\mathbf{a}'\mathbf{D}\mathbf{a} = \text{Tr}(\mathbf{a}'\mathbf{D}\mathbf{a}) = \text{Tr}(\mathbf{D}\mathbf{a}\mathbf{a}')$ , the expected log-kernel can be written as,

$$(\frac{n}{2} + \alpha_e - 1) \log(\tau_e) - \tau_e \left( \gamma_e + \frac{(\mathbf{y} - \mathbf{X}E_{\beta,\mathbf{u}}(\beta) - \mathbf{Z}E_{\beta,\mathbf{u}}(\mathbf{u}))'(\mathbf{y} - \mathbf{X}E_{\beta,\mathbf{u}}(\beta) - \mathbf{Z}E_{\beta,\mathbf{u}}(\mathbf{u})) + \text{Tr} \left( \begin{pmatrix} \mathbf{X}'\mathbf{X} & \mathbf{X}'\mathbf{Z} \\ \mathbf{Z}'\mathbf{X} & \mathbf{Z}'\mathbf{Z} \end{pmatrix} \begin{pmatrix} \text{Var}(\beta) & \text{Var}(\beta, \mathbf{u}) \\ \text{Var}(\mathbf{u}, \beta) & \text{Var}(\mathbf{u}) \end{pmatrix} \right)}{2} \right)$$

- which indicates that the approximate posterior for  $\tau_e$  is

$$\text{Ga} \left( \frac{n}{2} + \alpha_e, \gamma_e + \frac{(\mathbf{y} - \mathbf{X}E_{\beta,\mathbf{u}}(\beta) - \mathbf{Z}E_{\beta,\mathbf{u}}(\mathbf{u}))'(\mathbf{y} - \mathbf{X}E_{\beta,\mathbf{u}}(\beta) - \mathbf{Z}E_{\beta,\mathbf{u}}(\mathbf{u})) + \text{Tr} \left( \begin{pmatrix} \mathbf{X}'\mathbf{X} & \mathbf{X}'\mathbf{Z} \\ \mathbf{Z}'\mathbf{X} & \mathbf{Z}'\mathbf{Z} \end{pmatrix} \begin{pmatrix} \text{Var}(\beta) & \text{Var}(\beta, \mathbf{u}) \\ \text{Var}(\mathbf{u}, \beta) & \text{Var}(\mathbf{u}) \end{pmatrix} \right)}{2} \right).$$

- For  $\tau_u$ , the kernel and log-kernel are respectively

$$\begin{aligned} \text{Kernel: } & \left(\frac{\tau_u}{2\pi}\right)^{\frac{q}{2}} e^{-\frac{\tau_u \mathbf{u}' \mathbf{K}^{-1} \mathbf{u}}{2}} \tau_u^{\alpha_u-1} e^{-\gamma_u \tau_u} \\ \text{Log-kernel: } & (q/2 + \alpha_u - 1) \log(\tau_u) - \tau_u(\gamma_u + \mathbf{u}'\mathbf{K}^{-1}\mathbf{u}/2) \end{aligned}$$

- The expected log-kernel  $E_{-\tau_u}$  (Log-kernel) is

$$\begin{aligned} &= (q/2 + \alpha_u - 1) \log(\tau_u) - \tau_u(\gamma_u + E_{\beta,\mathbf{u}}(\mathbf{u}'\mathbf{K}^{-1}\mathbf{u})/2) \\ &= (q/2 + \alpha_u - 1) \log(\tau_u) - \tau_u(\gamma_u \text{Tr}(\mathbf{K}^{-1}E_{\beta,\mathbf{u}}(\mathbf{u}\mathbf{u}'))/2) \\ &= (q/2 + \alpha_u - 1) \log(\tau_u) - \tau_u(\gamma_u + E_{\beta,\mathbf{u}}(\mathbf{u})'\mathbf{K}^{-1}E_{\beta,\mathbf{u}}(\mathbf{u})/2 + \text{Tr}(\mathbf{K}^{-1}Var(\mathbf{u}))/2) \end{aligned}$$

- which indicates that the approximate posterior for  $\tau_u$  is

$$\text{Ga}\left(\frac{q}{2} + \alpha_u, \gamma_u + \frac{E_{\beta,\mathbf{u}}(\mathbf{u})' \mathbf{K}^{-1} E_{\beta,\mathbf{u}}(\mathbf{u}) + \text{Tr}(\mathbf{K}^{-1} \text{Var}(\mathbf{u}))}{2}\right).$$

- For  $\beta, \mathbf{u}$ , the kernel and log-kernel are respectively

$$\begin{aligned} \text{Kernel: } &= e^{-\frac{\tau_e(\mathbf{y}-\mathbf{X}\beta-\mathbf{Z}\mathbf{u})'(\mathbf{y}-\mathbf{X}\beta-\mathbf{Z}\mathbf{u})}{2} - \frac{\tau_u \mathbf{u}' \mathbf{K}^{-1} \mathbf{u}}{2}} \\ &= e^{-\frac{\tau_e(\mathbf{y}-(\mathbf{x} \ z) \ (\beta) \ (\mathbf{u}))'(\mathbf{y}-(\mathbf{x} \ z) \ (\beta) \ (\mathbf{u}))}{2}} e^{-\frac{\tau_u (\beta' \ u') \left( \begin{smallmatrix} \mathbf{0}_{p \times p} & \mathbf{0}_{p \times q} \\ \mathbf{0}_{q \times p} & \mathbf{K}^{-1} \end{smallmatrix} \right) (\beta) \ (\mathbf{u})}{2}} \\ \text{Log-kernel: } &= -\frac{\tau_e(\mathbf{y}-(\mathbf{x} \ z) \ (\beta) \ (\mathbf{u}))'(\mathbf{y}-(\mathbf{x} \ z) \ (\beta) \ (\mathbf{u}))}{2} - \frac{\tau_u (\beta' \ u') \left( \begin{smallmatrix} \mathbf{0}_{p \times p} & \mathbf{0}_{p \times q} \\ \mathbf{0}_{q \times p} & \mathbf{K}^{-1} \end{smallmatrix} \right) (\beta) \ (\mathbf{u})}{2} \end{aligned}$$

- The expected log-kernel  $E_{-\beta,\mathbf{u}}$ (Log-kernel) is

$$\begin{aligned} &= -\frac{E_{\tau_e}(\tau_e)(\mathbf{y}-(\mathbf{x} \ z) \ (\beta) \ (\mathbf{u}))'(\mathbf{y}-(\mathbf{x} \ z) \ (\beta) \ (\mathbf{u}))}{2} - \frac{E_{\tau_u}(\tau_u) (\beta' \ u') \left( \begin{smallmatrix} \mathbf{0}_{p \times p} & \mathbf{0}_{p \times q} \\ \mathbf{0}_{q \times p} & \mathbf{K}^{-1} \end{smallmatrix} \right) (\beta) \ (\mathbf{u})}{2} \\ &\propto -\frac{E_{\tau_e}(\tau_e) (\beta' \ u') \left( \begin{smallmatrix} \mathbf{x}' \mathbf{x} & \mathbf{x}' \mathbf{z} \\ \mathbf{z}' \mathbf{x} & \mathbf{z}' \mathbf{z} \end{smallmatrix} \right) (\beta) \ (\mathbf{u})}{2} + E_{\tau_e}(\tau_e) (\beta' \ u') \left( \begin{smallmatrix} \mathbf{X}' \mathbf{y} \\ \mathbf{Z}' \mathbf{y} \end{smallmatrix} \right) - \frac{E_{\tau_u}(\tau_u) (\beta' \ u') \left( \begin{smallmatrix} \mathbf{0}_{p \times p} & \mathbf{0}_{p \times q} \\ \mathbf{0}_{q \times p} & \mathbf{K}^{-1} \end{smallmatrix} \right) (\beta) \ (\mathbf{u})}{2} \\ &= -\frac{(\beta' \ u') \left( \begin{smallmatrix} E_{\tau_e}(\tau_e) \mathbf{x}' \mathbf{x} & E_{\tau_e}(\tau_e) \mathbf{x}' \mathbf{z} \\ E_{\tau_e}(\tau_e) \mathbf{z}' \mathbf{x} & E_{\tau_e}(\tau_e) \mathbf{z}' \mathbf{z} + E_{\tau_u}(\tau_u) \mathbf{K}^{-1} \end{smallmatrix} \right) (\beta) \ (\mathbf{u})}{2} + E_{\tau_e}(\tau_e) (\beta' \ u') \left( \begin{smallmatrix} \mathbf{X}' \mathbf{y} \\ \mathbf{Z}' \mathbf{y} \end{smallmatrix} \right) \end{aligned}$$

- which indicates that the approximate posterior for  $\beta, \mathbf{u}$  is

$$\mathcal{N}\left(E_{\tau_e}(\tau_e) \begin{pmatrix} E_{\tau_e}(\tau_e) \mathbf{X}' \mathbf{X} & E_{\tau_e}(\tau_e) \mathbf{X}' \mathbf{Z} \\ E_{\tau_e}(\tau_e) \mathbf{Z}' \mathbf{X} & E_{\tau_e}(\tau_e) \mathbf{Z}' \mathbf{Z} + E_{\tau_u}(\tau_u) \mathbf{K}^{-1} \end{pmatrix}^{-1} \begin{pmatrix} \mathbf{X}' \mathbf{y} \\ \mathbf{Z}' \mathbf{y} \end{pmatrix}, \begin{pmatrix} E_{\tau_e}(\tau_e) \mathbf{X}' \mathbf{X} & E_{\tau_e}(\tau_e) \mathbf{X}' \mathbf{Z} \\ E_{\tau_e}(\tau_e) \mathbf{Z}' \mathbf{X} & E_{\tau_e}(\tau_e) \mathbf{Z}' \mathbf{Z} + E_{\tau_u}(\tau_u) \mathbf{K}^{-1} \end{pmatrix}^{-1}\right).$$

```
#Arguments are
#epsilon: accuracy cut-off.
#iter: no of iterations
#Kinv: inverse of K, where p(u) = N(0, sigma^2_u K)
#Z: Predictor matrix for random effects
#X: Predictor matrix for fixed effects
#y: response vector
#taue_0: initial guess for residual precision.
#tauu_0: initial guess for random effect precision
#a.u, g.u: hyper-parameters of gamma prior for tauu
#a.e, g.e: hyper-parameters of gamma prior for taue

#Output are final estimates, plus iteration number when convergence was reached.
VB.mm<-function(epsilon,iter,Kinv,Z,X,y,taue_0,tauu_0,u0,beta0,a.e,g.e,a.u,g.u){
  n<-dim(X) [1]
  p<-dim(X) [2]
  q<-dim(Z) [2]
  W <-cbind(X,Z)
  WTW<-crossprod(W)
  WTY<-crossprod(W,y)
```

```

Kinval1<-matrix(0,p+q,p+q)
Kinval1[-c(1:p),-c(1:p)]<-Kinv

for(i in 1:iter){
  Vub <-solve(taue_0*WTW+tauu_0*Kinval1) #update Var(b,u)
  ub <-taue_0*Vub%*%WTY #update E(b,u)
  TrKinvub <- sum(diag(Kinval1%*%Vub))
  uKinvub <- t(ub)%*%Kinval1%*%ub
  tauu <- (a.u+0.5*q)/(g.u+0.5*as.numeric(uKinvub)+0.5*TrKinvub)
  tauu <- as.numeric(tauu)
  err <- y - W%*%ub
  TrWTWub <- sum(diag(WTW%*%Vub))
  taue <- (a.e+0.5*n)/(g.e+0.5*sum(err^2)+0.5*TrWTWub)
  taue <- as.numeric(taue)

  if(i > 1){
    diffub <- sqrt((ub-ub0)^2)/(abs(ub)+0.01)
    diffte <- abs(taue_0-taue)/(taue+0.01)
    difftu <- abs(tauu_0-tauu)/(tauu+0.01)
    diffvub <- sqrt((diag(Vub0) - diag(Vub))^2)/(diag(Vub))
    diff.all<-c(diffub,diffte,difftu,diffvub)
    if(max(diff.all) < epsilon) break
  }
  Vub0 <- Vub;ub0<-ub;taue_0<-taue;tauu_0<-tauu
  #Calculate relative change.
}

taue.param<-c((a.e+0.5*n),(g.e+0.5*sum(err^2)+0.5*TrWTWub))
tauu.param<-c((a.u+0.5*q),(g.u+0.5*uKinvub+0.5*TrKinvub))
param<-list(ub,Vub,taue.param,tauu.param,i)
names(param)<-c('betau_mean','betau_var','tau_e','tau_u','iter')
return(param)
}

```

The R code used to implement Gibbs sampling for this mixed model in lecture 18 is given below:

```

#Arguments are
#iter: no of iterations
#Z: Predictor matrix for random effects
#X: Predictor matrix for fixed effects
#y: response vector
#burnin: number of initial iterations to discard.
#taue_0: initial guess for residual precision.
#tauu_0: initial guess for random effect precision
#Kinu: inverse of K, where p(u) = N(0, \sigma^2_u K)
#a.u, b.u: hyper-parameters of gamma prior for tauu
#a.e, b.e: hyper-parameters of gamma prior for taue

normalmm.Gibbs<-function(iter,Z,X,y,burnin,taue_0,tauu_0,Kinv,a.u,b.u,a.e,b.e){
  n <-length(y) #no. observations
  p <-dim(X)[2] #no of fixed effect predictors.
  q <-dim(Z)[2] #no of random effect levels
  tauu<-tauu_0
  taue<-taue_0

```

```

beta0<-rnorm(p)
u0    <-rnorm(q,0, sd=1/sqrt(tauu))

#Building combined predictor matrix.
W<-cbind(X,Z)
WTW <-crossprod(W)
WTy <-crossprod(W,y)
library(mvtnorm)

#storing results.
par <-matrix(0,iter,p+q+2)
#Calculating log predictive densities
lppd<-matrix(0,iter,n)

#Create modified identity matrix for joint posterior.
I0 <-diag(p+q)
diag(I0)[1:p]<-0
I0[-c(1:p),-c(1:p)] <-Kinv

for(i in 1:iter){
#Conditional posteriors.
  uKinvu <- t(u0) %*% Kinv %*% u0
  uKinvu <-as.numeric(uKinvu)
  tauu <-rgamma(1,a.u+0.5*q,b.u+0.5*uKinvu)
#Updating component of normal posterior for beta,u
  Prec <-WTW + tauu*I0/taue
  P.mean<- solve(Prec)%*%WTy
  P.var <-solve(Prec)/taue
  betau <-rmvnorm(1,mean=P.mean,sigma=P.var)
  betau <-as.numeric(betau)
  err   <- y-W%*%betau
  taue <-rgamma(1,a.e+0.5*n,b.e+0.5*sum(err^2))
#storing iterations.
  par[i,]<-c(betau,1/sqrt(tauu),1/sqrt(taue))
  beta0 <-betau[1:p]
  u0    <-betau[p+1:q]
  lppd[i,]= dnorm(y,mean=as.numeric(W%*%betau),sd=1/sqrt(taue))
}

lppd      = lppd[-c(1:burnin),]
lppdest   = sum(log(colMeans(lppd)))           #Estimating lppd for whole dataset.
pwaic2    = sum(apply(log(lppd),2,FUN=var)) #Estimating effective number of parameters.
par <-par[-c(1:burnin),]
colnames(par)<-c(paste('beta',1:p,sep=''),paste('u',1:q,sep=''),'sigma_b','sigma_e')
mresult<-list(par,lppdest,pwaic2)
names(mresult)<-c('par','lppd','pwaic')
return(mresult)
}

```

Any other code you may use can be modified from the code given in Lecture 18. Possible things you may want to check include:

- convergence of Gibbs sampler
- comparing the empirical and approximate distributions using density plots.

```

#data<-read.csv('farmdata.csv')
data<-read.csv(file.choose())
#K   <-read.csv('Kmat.csv')
K<-read.csv(file.choose())
K   <-as.matrix(K)
Kinv<-solve(K)

n<-dim(data)[1]
q<-dim(Kinv)[1]
X<-table(1:n,data$flock) #flock is fixed effect
#Indicator matrix for parents.
Z2<-table(1:n,data$sire)
Z3<-cbind(Z2,table(1:n,data$dam))

```

## Importing and formatting the data

```

system.time(chain1<-normalmm.Gibbs(iterator=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=5,tauu_0=0.2,Kinv=Kinv))

running the Gibbs sampler, checking convergence, and calculating effective sample size

## Warning: package 'mvtnorm' was built under R version 4.0.5

##       user   system elapsed
##      4.68     0.17    4.87

system.time(chain2<-normalmm.Gibbs(iterator=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=1,tauu_0=1,Kinv=Kinv)

##       user   system elapsed
##      4.56     0.06    4.63

system.time(chain3<-normalmm.Gibbs(iterator=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=0.2,tauu_0=3,Kinv=Kinv)

##       user   system elapsed
##      4.47     0.05    4.52

library(coda)

## Warning: package 'coda' was built under R version 4.0.5

rml1<-as.mcmc.list(as.mcmc((chain1$par[1:4000,])))
rml2<-as.mcmc.list(as.mcmc((chain2$par[1:4000,])))
rml3<-as.mcmc.list(as.mcmc((chain3$par[1:4000,])))
rml4<-as.mcmc.list(as.mcmc((chain1$par[4000+1:4000,])))
rml5<-as.mcmc.list(as.mcmc((chain2$par[4000+1:4000,])))
rml6<-as.mcmc.list(as.mcmc((chain3$par[4000+1:4000,])))
rml<-c(rml1,rml2,rml3,rml4,rml5,rml6)

#Gelman-Rubin diagnostic.
gelman.diag(rml)[[1]]


##          Point est. Upper C.I.
## beta1      1.0001886  1.000618
## beta2      1.0003186  1.000664
## u1         1.0002849  1.000427
## u2         1.0004086  1.000672

```

```

## u3      1.0002609  1.000783
## u4      1.0002456  1.000754
## u5      1.0001370  1.000347
## u6      1.0001749  1.000684
## u7      1.0001805  1.000262
## u8      1.0000731  1.000364
## u9      1.0001241  1.000346
## u10     1.0000389  1.000371
## u11     0.9999831  1.000121
## u12     1.0000541  1.000331
## u13     1.0000604  1.000143
## sigma_b 1.0009851  1.002606
## sigma_e  1.0012429  1.002924

#effective sample size.
effectiveSize(rml)

##      beta1      beta2      u1      u2      u3      u4      u5      u6
## 24000.000 24000.000 24000.000 21692.020 23909.883 24000.000 23848.486 24000.000
##      u7      u8      u9      u10     u11     u12     u13 sigma_b
## 24000.000 25176.790 23691.012 23376.080 24000.000 24000.000 24000.000 10917.902
## sigma_e
## 8676.271

```

```
system.time(test1<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=0.2,tauu_0=0.2,u0=rn
```

### Estimating parameters of variational Bayes approximations

```

##    user  system elapsed
##    0.10    0.01    0.11

system.time(test2<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=5,tauu_0=1,u0=rnorm(1000,0,1))

##    user  system elapsed
##    0      0      0

system.time(test3<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=1,tauu_0=5,u0=rnorm(1000,0,1))

##    user  system elapsed
##    0      0      0

test1$iter

## [1] 20

test2$iter

## [1] 16

test3$iter

## [1] 18

```

```
#Comparing point estimates/posterior means
```

```
chain.all<-rbind(chain1$par,chain2$par,chain3$par)
```

```

#beta, u
cbind(test1$betau_mean,test2$betau_mean,test3$betau_mean,colMeans(chain.all[,1:15]))
```

Comparing estimates obtained from Gibbs sampling and Variational Bayes

```

##      [,1]      [,2]      [,3]      [,4]
## 1  4.15443108 4.1544310 4.15443105 4.15312045
## 2  4.70775387 4.7077539 4.70775388 4.70932202
## 1  0.81178839 0.8117883 0.81178835 0.80841251
## 2 -1.09477605 -1.0947760 -1.09477600 -1.09169227
## 3  0.21037808 0.2103781 0.21037809 0.21098242
## 4  0.47799077 0.4779907 0.47799075 0.47601605
## 5 -0.53525688 -0.5352568 -0.53525681 -0.53454094
## 6 -0.70912449 -0.7091244 -0.70912442 -0.70499465
## 7  0.51809256 0.5180925 0.51809251 0.51611965
## 8 -0.16314311 -0.1631431 -0.16314311 -0.16250230
## 9 -0.14790542 -0.1479054 -0.14790543 -0.14939405
## 10 0.13369963 0.1336996 0.13369961 0.13095244
## 11 0.15439310 0.1543930 0.15439306 0.15130713
## 12 -0.07566564 -0.0756656 -0.07566562 -0.07495579
## 13 0.57227999 0.5722800 0.57227999 0.57072883
```

*#Variances.*

```

sigmae2<-sigmau2<-rep(0,3)
sigmae2[1]<-test1$tau_e[2]/(test1$tau_e[1]-1)
sigmau2[1]<-test1$tau_u[2]/(test1$tau_u[1]-1)
sigmae2[2]<-test2$tau_e[2]/(test2$tau_e[1]-1)
sigmau2[2]<-test2$tau_u[2]/(test2$tau_u[1]-1)
sigmae2[3]<-test3$tau_e[2]/(test3$tau_e[1]-1)
sigmau2[3]<-test3$tau_u[2]/(test3$tau_u[1]-1)
cbind(sigmau2,mean(chain.all[,16]^2))
```

```

##      sigmau2
## [1,] 0.4304444 0.4649196
## [2,] 0.4304445 0.4649196
## [3,] 0.4304444 0.4649196
cbind(sigmae2,mean(chain.all[,17]^2))
```

```

##      sigmae2
## [1,] 0.04876863 0.05242605
## [2,] 0.04876872 0.05242605
## [3,] 0.04876868 0.05242605
```

*#Comparing posterior distributions.*

```

par(mfrow=c(5,4))
mlim<-quantile(chain.all[,1],c(0.005,0.995))
curve(dnorm(x,mean=test1$betau_mean[1],sd=sqrt(test1$betau_var[1,1])),ylab='Density',main='',xlim=mlim,
lines(density(chain.all[,1]))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)

mlim<-quantile(chain.all[,2],c(0.005,0.995))
curve(dnorm(x,mean=test1$betau_mean[2],sd=sqrt(test1$betau_var[2,2])),ylab='Density',main='',xlim=mlim,
lines(density(chain.all[,2]))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)
```

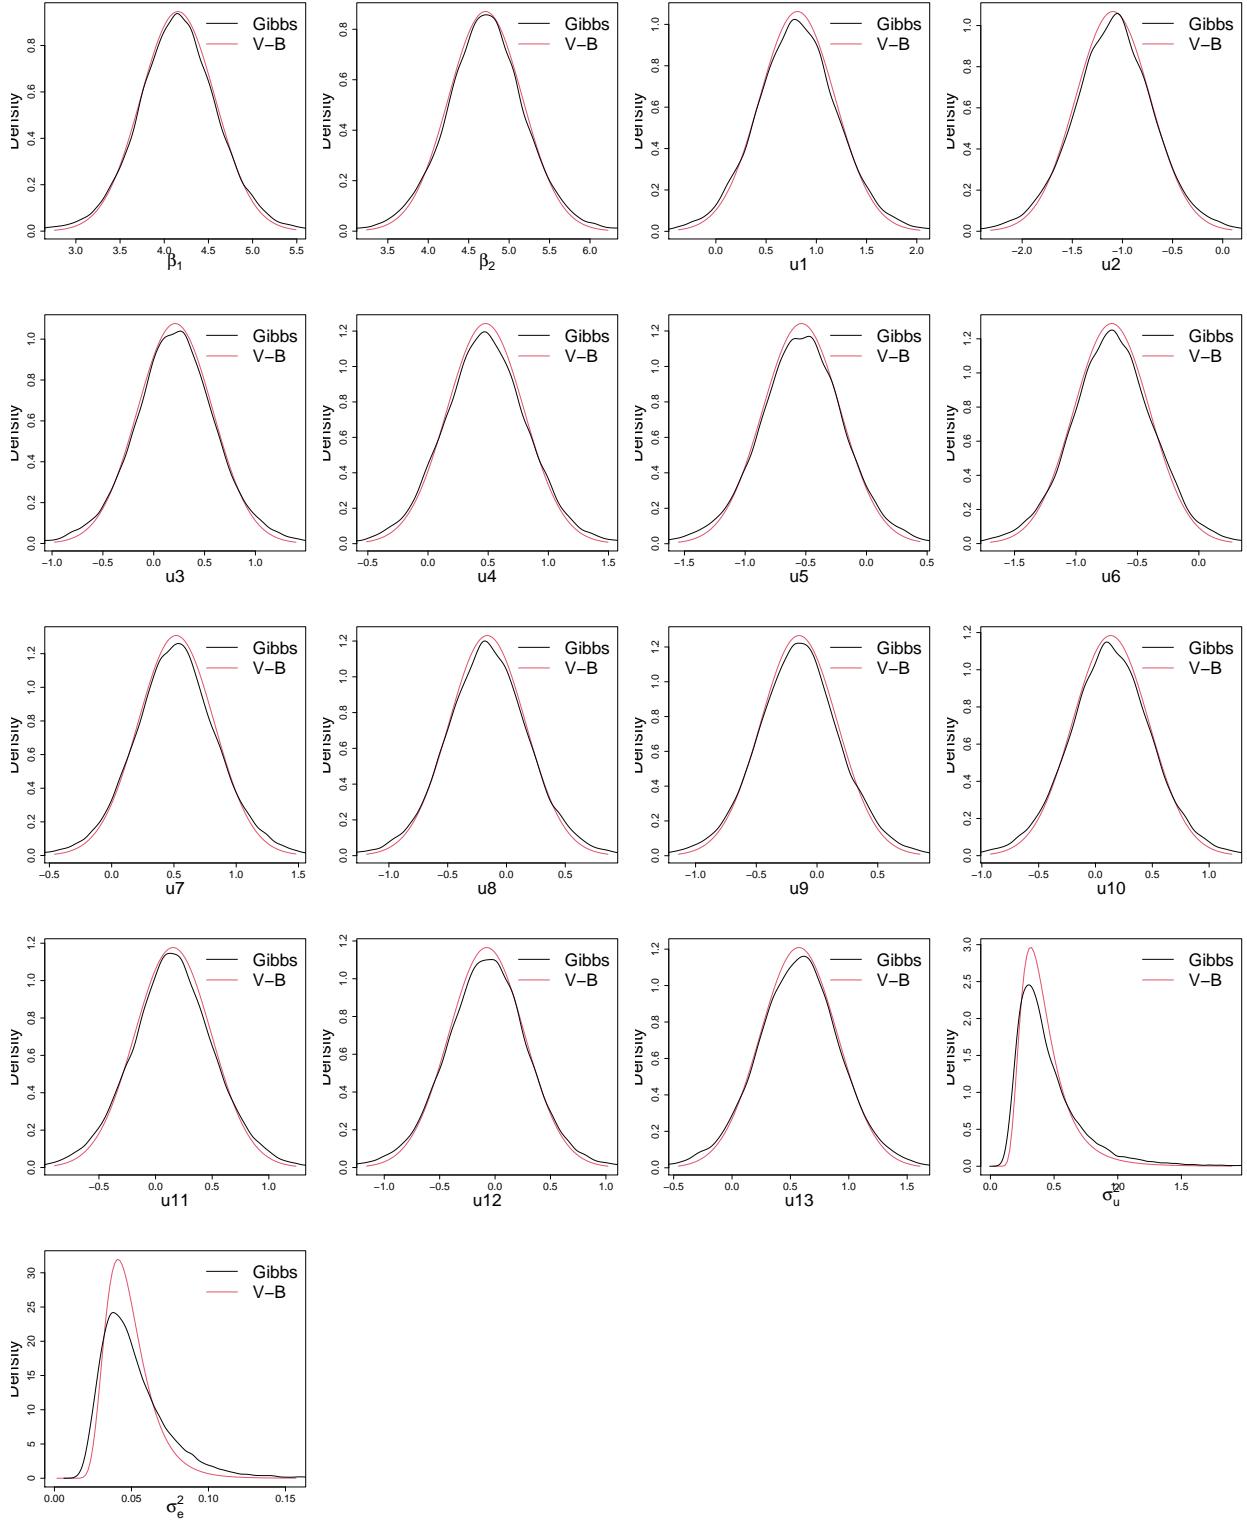
```

#Repeat for random effects.
for(i in 1:13){
  mlim<-quantile(chain.all[,i+2],c(0.005,0.995))
  curve(dnorm(x,mean=test1$betau_mean[i+2],sd=sqrt(test1$betau_var[i+2,i+2])),ylab='Density',main='',xlim=c(0,mlim))
  lines(density(chain.all[,i+2]))
  legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)
}

mlim<-quantile(chain.all[,16]^2,c(0.005,0.995))
curve(dgamma(1/x,shape=test1$tau_u[1],rate=test1$tau_u[2])*x^(-2),ylab='Density',main='',xlim=c(0,mlim))
lines(density(chain.all[,16]^2))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)

mlim<-quantile(chain.all[,17]^2,c(0.005,0.995))
curve(dgamma(1/x,shape=test1$tau_e[1],rate=test1$tau_e[2])*x^(-2),ylab='Density',main='',xlim=c(0,mlim))
lines(density(chain.all[,17]^2))
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5)

```



By determining the joint approximate posterior for  $\beta, \mathbf{u}$  rather than separate independent approximate posteriors for  $\beta, \mathbf{u}$  as in lecture 18, we improve the accuracy of approximate inference. This is because  $\mathbf{X}$  and  $\mathbf{Z}$  are not independent, and so contracting independent approximate posteriors for  $\beta$  and  $\mathbf{u}$ , we ignore the information contained in  $\mathbf{X}'\mathbf{Z}$ .